

Heart Disease Prediction

Name:Saraansh Chikara

```
In [4]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [6]: df=pd.read_csv("Heart.csv")
```

```
In [7]: df.head()
```

Out[7]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [53]: df.isnull().sum()
```

```
Out[53]: age          0  
sex          0  
cp           0  
trestbps     0  
chol         0  
fbs          0  
restecg      0  
thalach      0  
exang        0  
oldpeak      0  
slope        0  
ca           0  
thal         0  
target       0  
dtype: int64
```

```
In [112]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null   int64
 1   sex         303 non-null   int64
 2   cp          303 non-null   int64
 3   trestbps    303 non-null   int64
 4   chol        303 non-null   int64
 5   fbs         303 non-null   int64
 6   restecg     303 non-null   int64
 7   thalach     303 non-null   int64
 8   exang       303 non-null   int64
 9   oldpeak     303 non-null   float64
10   slope       303 non-null   int64
11   ca          303 non-null   int64
12   thal        303 non-null   int64
13   target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
In [113]: df.describe()
```

```
Out[113]:
```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733	1.039604
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794	1.161075
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000	0.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000	0.000000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000	0.800000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000	1.600000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000	6.200000

```
In [27]: a=df.shape  
print(f"The data set has {a[0]} rows and {a[1]} columns")
```


The data set has 303 rows and 14 columns

Data Visualisation


```
In [67]: plt.hist(df["age"], bins=30, edgecolor="Black")  
plt.title('Histogram of Age')  
plt.xlabel('Age')  
plt.ylabel('Frequency')  
plt.show()
```



```
In [8]: df.boxplot(column=['trestbps'])  
plt.title('Boxplot of Resting Blood Pressure')  
plt.show()
```



```
In [66]: correlation_matrix = df.corr()
plt.figure(figsize=(20, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=.25)
plt.title('Correlation Matrix')
plt.show()
```



Data Modelling

1) Logistic Regression

```
In [9]: y=df['target']  
x=df.drop(columns='target')
```

```
In [11]: from sklearn.model_selection import train_test_split  
x_train, x_test,y_train, y_test=train_test_split(x,y, test_size=0.30, random_state=40)
```

```
In [12]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()  
x_train = sc.fit_transform(x_train)  
x_test = sc.transform(x_test)
```

```
In [14]: from sklearn.linear_model import LogisticRegression  
model=LogisticRegression()
```

```
In [15]: model.fit(x_train, y_train)
```

```
Out[15]: 

▼ LogisticRegression



LogisticRegression()


```

```
In [16]: model.intercept_
```

```
Out[16]: array([0.09066276])
```



```
In [17]: model.coef_
```

```
Out[17]: array([[ 0.00210412, -0.75695699,  0.8695811 , -0.41305306, -0.1693742 ,
                -0.0412552 ,  0.38040936,  0.5085677 , -0.27251041, -0.67049861,
                0.28262575, -0.59555475, -0.38458101]])
```

```
In [18]: y_pred=model.predict(x_train)
print(y_pred)
```

```
[1 1 0 0 0 0 0 1 1 0 1 1 1 0 0 0 1 1 1 1 1 1 1 0 0 0 0 1 1 1 1 0 0 0 0 1
 0 1 1 1 1 1 0 1 1 1 1 1 1 0 0 1 1 1 0 0 1 1 0 1 0 0 1 1 1 0 1 1 1 1 1 1
 1 1 1 1 0 1 1 1 1 1 0 1 0 0 0 1 0 1 1 0 0 1 1 0 0 1 1 0 1 1 1 0 1 1 0 0 1
 1 1 0 1 1 1 1 1 1 1 0 1 0 1 0 0 1 1 0 1 1 1 0 0 1 0 1 1 0 0 0 1 0 1 1 0 1
 0 1 1 1 0 1 1 0 0 0 1 0 1 0 1 1 1 0 0 1 1 1 0 1 0 0 1 0 1 1 1 0 1 0 1 1 1
 1 1 0 0 1 1 1 0 0 1 0 0 0 1 0 0 1 1 1 1 1 1 1 0 0 0 1 0]
```

```
In [19]: y_prediction=model.predict(x_test)
print(y_prediction)
```

```
[1 1 1 0 1 1 1 0 0 1 0 1 1 0 0 1 1 1 0 0 1 1 1 1 0 1 0 1 1 1 1 0 0 1 1 1 0
 0 0 1 1 1 1 0 1 1 1 0 0 0 1 1 1 0 1 1 0 1 0 0 1 1 1 1 1 1 1 0 0 0 0 0 0 0
 1 0 1 0 0 1 0 1 0 1 0 1 1 1 0 1 1]
```

```
In [21]: from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
cm=confusion_matrix(y_test, y_prediction)
cm
```

```
Out[21]: array([[35,  5],
                [ 2, 49]], dtype=int64)
```

```
In [22]: accuracy=accuracy_score(y_test, y_prediction)
print(f"Accuracy score is :{accuracy}")
```


Accuracy score is :0.9230769230769231

```
In [23]: train_score=model.score(x_train,y_train)
test_score=model.score(x_test,y_test)
print(f"Training score is :{train_score}")
print(f"Testing score is :{test_score}")
```

Training score is :0.8349056603773585

Testing score is :0.9230769230769231

```
In [24]: # Plot the confusion matrix as a heatmap
sns.heatmap(cm, annot=True, fmt=".2f", cmap="viridis")
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
In [25]: print(f"Model score :{model.score(x_test,y_test)}")
```

Model score :0.9230769230769231

```
In [28]: from sklearn.model_selection import cross_val_score
accuracies = cross_val_score(estimator = model, X = x_train, y = y_train, cv = 5)
print(accuracies)
```

[0.74418605 0.76744186 0.80952381 0.88095238 0.73809524]

```
In [29]: print("Accuracy mean: {:.2f} %".format(accuracies.mean()*100))
print("Accuracy max: {:.2f} %".format(accuracies.max()*100))
```

Accuracy mean: 78.80 %

Accuracy max: 88.10 %

```
In [65]: a=accuracy_score(y_test, y_prediction)
p=precision_score(y_test, y_prediction)
r=recall_score(y_test, y_prediction)
f1=f1_score(y_test, y_prediction)
print(f"Accuracy score is :{a}")
print(f"Precision score is:{p}")
print(f"Recall score is :{r}")
print(f"F1 score is :{f1}")
```

Accuracy score is :0.9230769230769231

Precision score is:0.9074074074074074

Recall score is :0.9607843137254902

F1 score is :0.9333333333333333

Decision Tree

```
In [31]: from sklearn.model_selection import train_test_split  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.metrics import accuracy_score
```

```
In [32]: X = df.drop(columns='target')  
Y = df['target']
```

```
In [33]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, random_state=42)
```

```
In [34]: from sklearn.preprocessing import StandardScaler  
sc = StandardScaler()  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)
```

```
In [35]: clf = DecisionTreeClassifier()
```

```
In [36]: clf.fit(X_train, Y_train)
```

```
Out[36]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [49]: Y_pred=model.predict(X_train)
print(Y_pred)
```

```
[1 1 1 1 1 1 0 1 0 0 1 1 0 1 1 1 1 0 1 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 1 0
 0 0 1 1 0 0 1 1 0 1 0 0 0 0 1 1 1 1 0 0 0 0 1 1 1 1 0 1 1 0 1 0 0 1 0 1 0
 1 1 1 1 1 0 1 0 0 0 0 0 1 1 1 1 1 1 1 0 1 0 0 1 1 1 1 1 1 0 0 0 0 1 1 1
 1 1 0 1 1 1 0 1 0 1 0 1 1 1 0 1 1 0 1 1 1 1 1 1 1 0 0 1 1 0 0 1 0 0 0
 0 1 1 1 0 0 1 0 1 1 0 1 0 0 0 1 1 1 1 1 1 0 1 0 0 0 1 1 1 0 0 1 0 0 1 1 0
 1 0 0 0 1 1 0 0 1 1 0 0 0 1 1 1 1 1 0 1 1 0 1 1 1 0 1]
```

```
In [42]: Y_prediction = clf.predict(X_test)
print(Y_prediction)
```

```
[0 1 1 0 1 1 1 0 0 0 1 0 1 0 1 1 1 0 0 0 1 0 1 0 0 0 1 1 0 1 0 0 0 0 1 0 0
 1 0 1 1 1 1 0 0 0 1 1 1 0 0 0 1 1 0 0 1 1 0 0 0 1 0 0 1 0 0 1 1 1 1 1 1 1
 0 1 1 0 0 0 1 0 1 0 0 1 0 0 0 0 1]
```


```
In [52]: train2_score=model.score(X_train,Y_train)
test2_score=model.score(X_test,Y_test)
print(f"Training score is :{train2_score}")
print(f"Testing score is :{test2_score}")
```

```
Training score is :0.8537735849056604
Testing score is :0.8461538461538461
```

```
In [56]: from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
cm=confusion_matrix(Y_test, Y_prediction)
cm
```

```
Out[56]: array([[31, 10],
               [17, 33]], dtype=int64)
```

```
In [54]: sns.heatmap(cm, annot=True, fmt=".2f", cmap="viridis")  
plt.title('Confusion Matrix')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.show()
```



```
In [55]: print(f"Model score :{model.score(X_test,Y_test)}")
```

Model score :0.8461538461538461

```
In [63]: A=accuracy_score(Y_test, Y_prediction)
P=precision_score(Y_test, Y_prediction)
R=recall_score(Y_test, Y_prediction)
F1=f1_score(Y_test, Y_prediction)
print(f"Accuracy score is :{A}")
print(f"Precision score is:{P}")
print(f"Recall score is :{R}")
print(f"F1 score is :{F1}")
```

Accuracy score is :0.7032967032967034

Precision score is:0.7674418604651163

Recall score is :0.66

F1 score is :0.7096774193548386

```
In [ ]:
```